



Carisurg

Assignment 20

Handover Documentation

Due Date: July 25th, 2026.

Prepared by **Josiah-John Green**

Project Summary

This repo trains and evaluates a machine-learning classifier that predicts a patient's Emergency Severity Index (ESI 1–5) from ten triage-time features (age, vitals, chief-complaint flags, arrival mode).

It was built on a public Yale EMLLC dataset (55,121 encounters) as a proof of feasibility while the Phase 1 Caribbean data-capture tool is developed in parallel; it **has not** been validated on Mercer General data and is not cleared for clinical use.

The Final Model Decision

The model used is **Logistic Regression**, ``class_weight="balanced"`,` pinned in ``config.yaml``.

Why?

Across six models tried in Weeks 6–7 (Dummy, Decision Tree, Logistic Regression, Random Forest, tuned Random Forest, Gradient Boosting), Logistic Regression has the lowest accuracy but by far the highest recall on ESI-1 — the most critical patient class (0.750, versus 0.000 for Random Forest and 0.375 for Gradient Boosting).

For a triage system, missing critical patients is a worse failure than a lower overall accuracy score.

- Full **comparison**: ``modules/week-08/docs/week08_selection.md``
- Full **reasoning**: ``modules/week-08/docs/decisions/week08_choice.md``

How to Run?

```
```bash

git clone https://github.com/j0hnc0d3s/carisurg.git

cd carisurg/modules/week-08

python3 -m venv venv

source venv/bin/activate

pip install -r requirements.txt

placed triage_cleaned_v1.csv in dataset/ (see "Where the Data Lives" below)

python scripts/train.py --config config.yaml

```
```

Expect console output ending in a metrics block (accuracy, ESI-1 recall, etc.) and three new files under `records/`: `final\_model.joblib`, `final\_scaler.joblib`, `final\_metrics.json`.

To confirm the pipeline itself is healthy before trusting any of that:

```
```bash

python -m pytest -v

```
```

Both tests should pass in a few seconds. They check the data schema and run a training smoke test on synthetic data — they do not check model quality.

Where the Data Lives?

1. **Source:** Yale Emergency Medicine Machine Learning Consortium (Yale EMMLC), a US academic hospital's de-identified ED triage export.
2. **File expected at:** `dataset/triage\_cleaned\_v1.csv` — git-ignored, not committed to this repo. See `modules/week-08/README.md` for provenance and how to obtain it.
3. **Access:** limited to CariSurg MedTech Pathways programme participants and tutors. Do not redistribute the CSV outside the programme, upload it to any third-party service, or commit it (or any copy/derivative of it, including model artefacts trained on it) to a public repository.
4. **Governance status:** external, de-identified, public research dataset — not Mercer General patient data, and no IRB/ethics approval has been sought or is required for it. De-identified does not mean ungoverned: the access restriction above still applies. Any future work using real Mercer General patient data will require separate ethics review before this pipeline (or any successor) touches it.

Known Limitations

1. ESI-1 recall is measured on only 16 test patients (77 of 55,121 encounters overall). Every recall number in ``modules/week-08/docs/week08-selection.md`` should be read as a consistent pattern across models, not a precise estimate — a single patient's outcome shifts the ESI-1 recall figure by more than 6 percentage points.
2. This is not Caribbean data. No result in this repo says anything about performance at Mercer General until validated against local ED data collected under Phase 1. Single-site data of any kind carries a real risk of distribution shift when applied elsewhere.
3. The pinned model uses only 10 features, and demographics are excluded by design, not oversight.
4. No race, ethnicity, or other demographic column is in the feature set — this was a deliberate fairness choice from the Week 5 memo, not a gap to fill in later. It does mean the model draws on far less than a clinician's full judgement, which caps how good it can realistically get.

Who to Ask?

| Question About | Contact |
|--|--|
| Model Choice | Josiah-John Green |
| Dataset Access | Martina Griffith (Clinical IT Lead) |
| Clinical Validity of ESI
Predictions/Triage Workflows | Dr Marcus Reyes (Consultant EP)
or Dr De Freitas (Clinical Sponsor) |
| Programme Process/Submission
Logistics | Carisurg Tutor Team, in ‘#ask-a-tutor’ |

Appendices

Supplementary Links

1. The GitHub link can be found [here](#)
2. The decision journal can be found in the docs for this week.

Glossary

A reference index for clinical, technical, and research terminology used throughout this proposal. Intended for any reviewer — clinical or non-clinical — who needs a quick definition without breaking flow.

Clinical & Vital Sign Terms

| Term | Meaning |
|------|--|
| ED | Emergency Department |
| ESI | Emergency Severity Index — the 5-level triage scale used at Mercer General (1 = immediate, 5 = non-urgent) |
| SBP | Systolic Blood Pressure — the peak pressure during a heartbeat (top number in a BP reading) |
| DBP | Diastolic Blood Pressure — the pressure between heartbeats (bottom number in a BP reading) |
| MAP | Mean Arterial Pressure — average pressure driving blood to organs across a cardiac cycle; calculated as $(SBP + 2 \times DBP) / 3$ |
| GCS | Glasgow Coma Scale — score from 3–15, measuring level of consciousness |
| RR | Respiratory Rate — breaths per minute |

| | |
|------|---|
| FiO2 | Fraction of Inspired Oxygen — the percentage of oxygen a patient is receiving (room air = 21%) |
| SpO2 | Oxygen Saturation — percentage of haemoglobin carrying oxygen, measured via pulse oximeter |
| SBAR | Situation, Background, Assessment, Recommendation — a structured format for clinical handover communication |

Healthcare Systems & Infrastructure Terms

| Term | Meaning |
|---------------|---|
| EHR | Electronic Health Record — digital patient record system; Mercer currently has none for triage, hence the Phase 1 proposal |
| LMIC | Low- and Middle-Income Country — a World Bank income classification used in global health literature |
| SIDS | Small Island Developing State — a UN classification for small island nations facing shared infrastructure constraints |
| Boarding time | The gap between a disposition decision (e.g. “admit”) and the patient physically leaving the ED, usually due to no ward bed being available |

AI / Machine Learning & Risk Terms

| Term | Meaning |
|-----------------------|--|
| AI | Artificial Intelligence |
| ML | Machine Learning |
| AUROC / AUC | Area Under the Receiver Operating Characteristic curve — a model performance metric from 0.5 to 1.0; higher is better |
| XGBoost | Extreme Gradient Boosting — a machine learning algorithm using an ensemble of decision trees, common for structured/tabular data |
| Rule-based classifier | A decision-making system using explicit if/then logic rather than a trained statistical model |
| PPV | Positive Predictive Value — the proportion of positive alerts that turn out to be true positives. Low PPV drives alert fatigue. |
| Automation bias | The tendency for a human user to over-trust or defer to an automated system's output, even when contradicted by their own judgment |
| Distribution shift | When a model's performance degrades because the data it encounters in deployment differs statistically from the data it was trained on |
| Human-in-the-loop | A system design principle requiring a human to review, confirm, or be able to override an AI-generated output before it takes effect |

Research & Documentation Terms

| Term | Meaning |
|-------------------|--|
| DOI | Digital Object Identifier — a permanent link format for academic papers |
| APA | American Psychological Association — the citation style used throughout this proposal (7th edition) |
| Scoping review | A literature review mapping the breadth of existing research, without necessarily assessing study quality in depth |
| Systematic review | A more rigorous literature review following a defined search and inclusion protocol |